Graph-CoT: from Neo4j optimization to vLLM re-entry

Reduce DB work; then measure when tool results re-enter the model queue.

A Cypher reduces traversal work

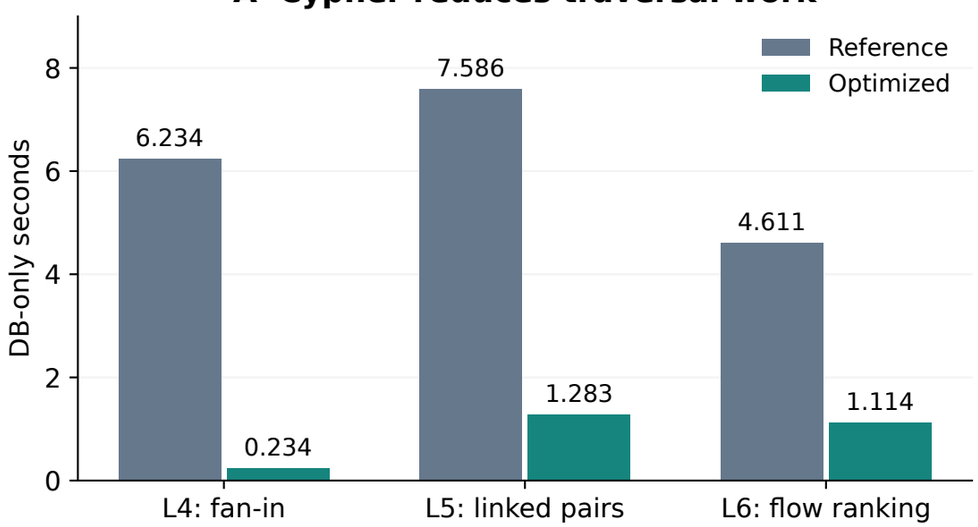

Historical Python study; 4 samples per tool/version.
Same query semantics; 256MiB page cache.

B Heap + cache: reliability and speed

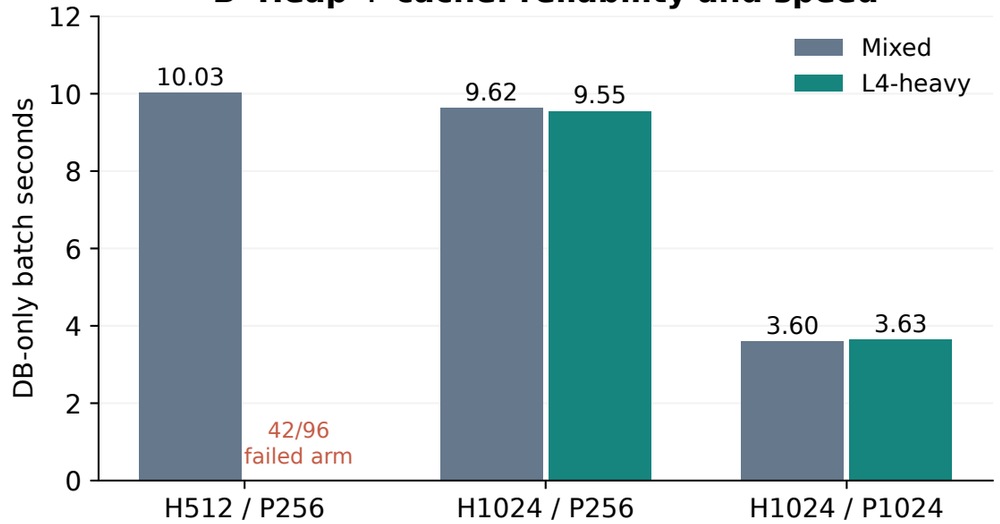

MiB; 3-block medians. Mixed: 240/240 per arm.
L4-heavy: failed baseline time omitted; others 96/96.

C DB gains do not guarantee E2E gains

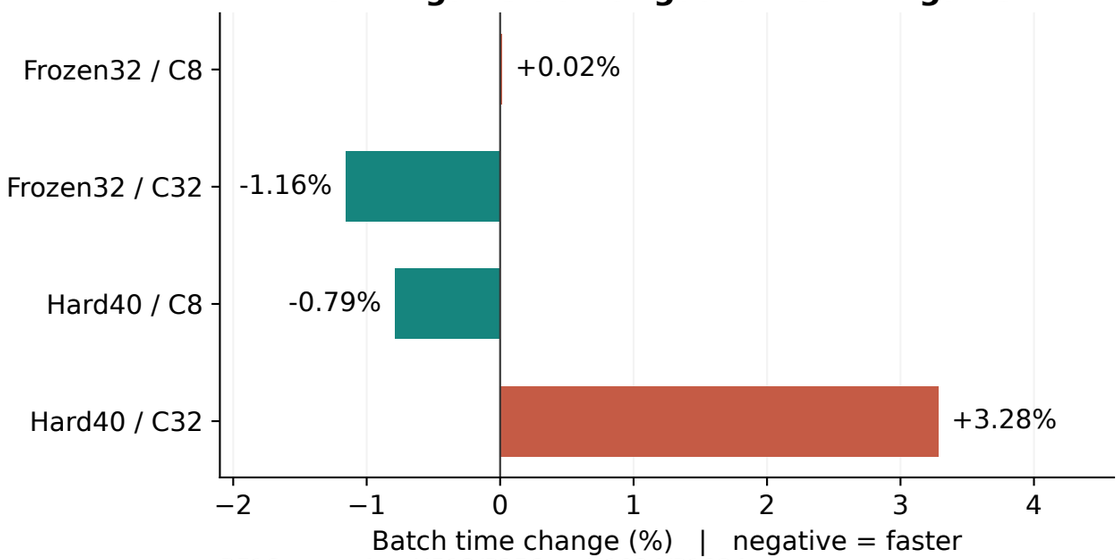

BF16 + neo4rs; H512/P256 vs H1024/P1024.
864/864 correct, 0 failed/dropped; 3-block medians.

D Same tail request, later model admission

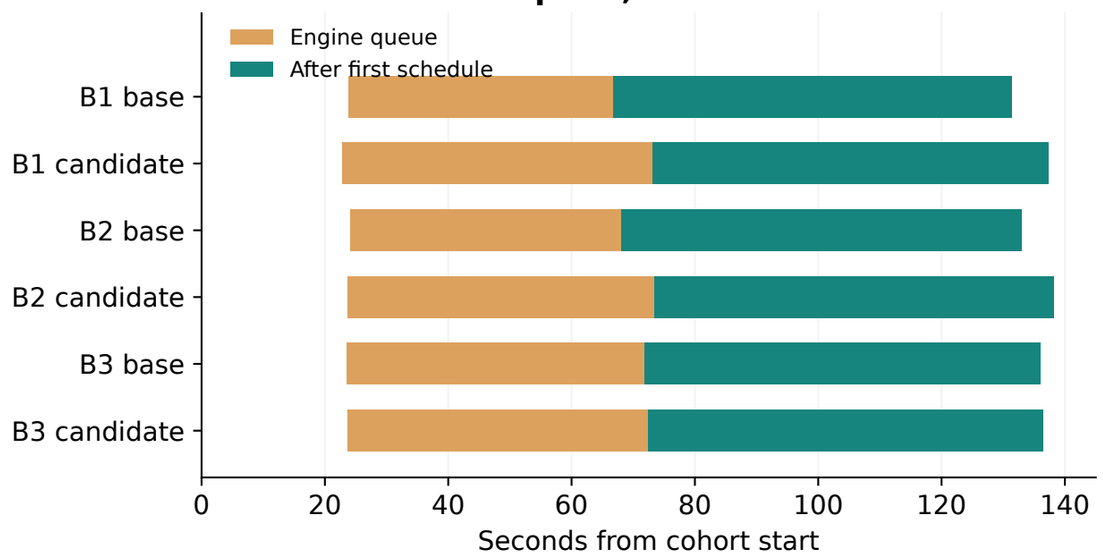

Hard40/C32: same request/evidence, 1,728 final tokens.
Blank prefix is earlier episode work, not idle time.

Separate studies, workloads and budgets; do not multiply gains. Finite cohorts on a shared host, not sustained capacity.
Source values and provenance: evidence/overview.json. Full interpretation: RESULTS.md